

Chapter 1

General-Purpose Block Preconditioners

In this document we will demonstrate how to use the general-purpose block preconditioners implemented in `oomph-lib`. This tutorial follows from the [Block Preconditioners](#) tutorial, which provides an overview of `oomph-lib`'s generic block preconditioning framework.

We use the `Problem` described in the [Bending of a Cantilever Beam](#) tutorial to illustrate the key concepts.

1.1 Introduction

In this section we define the four (distributed) general purpose block preconditioning methodologies. To recap, all `oomph-lib` problems are solved in a Newton iteration which requires the repeated solution of linear systems of the form

$$J\delta\mathbf{x} = -\mathbf{r}$$

where J is the Jacobian matrix, $\mathbf{r}$ is the vector of residuals and $\delta\mathbf{x}$ is the Newton correction. We divide the DOFs in the two-dimensional cantilever problem into two subsets corresponding to the x and y nodal positions.

$$\begin{bmatrix} J_{xx} & J_{xy} \\ J_{yx} & J_{yy} \end{bmatrix} \cdot \begin{bmatrix} \delta\mathbf{x} \\ \delta\mathbf{y} \end{bmatrix} = - \begin{bmatrix} \mathbf{r}_x \\ \mathbf{r}_y \end{bmatrix}$$

Utilising this partitioning we will describe four (distributed) general purpose block preconditioning methodologies. (Left) preconditioning represents a transformation of the original linear system to

$$P^{-1}J\delta\mathbf{x} = -P^{-1}\mathbf{r}$$

with the aim of accelerating the convergence of Krylov subspace iterative methods such as GMRES or CG. The application of the preconditioner requires the solution of

$$P\mathbf{z} = \mathbf{w}$$

for $\mathbf{z}$ at each Krylov iteration.

1.1.1 Block Diagonal Preconditioning

We drop the off-diagonal blocks to form the block diagonal preconditioner

$$P_{BD} = \begin{bmatrix} J_{xx} & \\ & J_{yy} \end{bmatrix}.$$

the application of this preconditioner requires the solution of the subsidiary systems J_{xx} and J_{yy} .

1.1.2 Block Diagonal Preconditioning with Two-Level Parallelisation

The two-subsystem systems in the block diagonal preconditioner (involving J_{xx} and J_{yy}) can be solved in any order. In a parallel computation we can either solve the two systems one after the other using the full set of processes for the solution of each linear system. An alternative is to solve all the subsidiary systems simultaneously, using only a subset of processes for each system. We refer to this technique as two-level parallelisation and note that this approach is particularly useful if the linear solvers do not have good parallel scaling properties.

1.1.3 Upper Block Triangular Preconditioning

An alternative to block diagonal preconditioning is block triangular preconditioning in which only off diagonal blocks on one side of the diagonal are dropped. For example, in the block-upper triangular preconditioner

$$P_{BUT} = \begin{bmatrix} J_{xx} & J_{xy} \\ & J_{yy} \end{bmatrix}$$

the block below the diagonal (J_{yx}) has been dropped. In addition to the two subsidiary solves for J_{xx} and J_{yy} this preconditioner requires a matrix-vector product involving J_{xy} .

1.1.4 Lower Block Triangular Preconditioning

Similarly we can define a lower triangular block preconditioner

$$P_{BLT} = \begin{bmatrix} J_{xx} & \\ J_{yx} & J_{yy} \end{bmatrix}.$$

1.2 Application

In this section we demonstrate the use of oomph-lib's general-purpose block preconditioners. All general-purpose block preconditioners are derived from the base class `GeneralPurposeBlockPreconditioner` (which is itself derived from the `BlockPreconditioner` class).

By default all general purpose block preconditioners use `SuperLUPreconditioner` as the preconditioner for the subsidiary systems (J_{xx} and J_{yy} in the [Introduction](#)). `SuperLUPreconditioner` is a wrapper to both the `SuperLU` direct solver and the `SuperLU Dist` distributed direct solver. Often we seek to replace this direct solver preconditioning with an inexact solver to make the preconditioner more efficient. To use an alternative subsidiary preconditioner we must define a function to return new instances of the chosen type of preconditioner (inexact solver). For example

```
//hypre_helper=====
/// The function get_hyre_preconditioner() returns an instance of
/// HyprePreconditioner to be used as a subsidiary preconditioner in a
/// GeneralPurposeBlockPreconditioner
//=====
namespace Hypre_Subsidary_Preconditioner_Helper
{
    Preconditioner* get_hyre_preconditioner()
    {
        return new HyprePreconditioner;
    }
} // end_of_hyre_helper
```

would return instances of `HyprePreconditioner`, a wrapper to the distributed `Hypre BoomerAMG` implementation of classical AMG. Later we will pass a pointer to this function to the block preconditioner to enable the use of `HyprePreconditioner` as a subsidiary preconditioner. Note that the function only creates the subsidiary preconditioner – it will be deleted automatically by the master preconditioner when it is no longer required.

The rest of the section is concerned with the main function, and in particular setting up the preconditioner for use.

```
//=====start_of_main=====
/// Driver for cantilever beam loaded by surface traction and/or
/// gravity
//=====
int main(int argc, char* argv[])
{
```

Given an instance of the problem,

```
//Set up the problem
CantileverProblem<MySolidElement<RefineableQPVDElement<2,3> > > > problem;
```

we specify GMRES as the linear solver. If available, we use the `TrilinosAztecOOSolver` wrapper to the `Trilinos AztecOO` implementation of GMRES. (This is the only distributed implementation of GMRES in oomph-lib.)

```
// use trilinos gmres if available
#ifdef OOMPH_HAS_TRILINOS
    TrilinosAztecOOSolver* solver_pt = new TrilinosAztecOOSolver;
    solver_pt->solver_type() = TrilinosAztecOOSolver::GMRES;
#else
    GMRES<CRDoubleMatrix>* solver_pt = new GMRES<CRDoubleMatrix>;
#endif
```

`GeneralPurposeBlockPreconditioner` is the base class for all general purpose block preconditioners.

```
// Pointer to general purpose block preconditioner base class
GeneralPurposeBlockPreconditioner<CRDoubleMatrix>* prec_pt = 0;
```

We introduced four general purpose block preconditioning methodologies in the [Introduction](#). The next step is to construct one of these preconditioners.

- **Block Diagonal Preconditioning.** This is implemented in the class `BlockDiagonalPreconditioner`.

```
// Standard Block Diagonal
prec_pt = new BlockDiagonalPreconditioner<CRDoubleMatrix>;
```

- **Enabling Two-Level Block Diagonal Preconditioning.** By default two-level preconditioning is disabled and hence `enable_two_level_parallelisation()` must have been called. Once this is done, each subsidiary system will be solved on an (as near to) equal size subset of processes.

```
// Two Level Block Diagonal
prec_pt = new BlockDiagonalPreconditioner<CRDoubleMatrix>;
dynamic_cast<BlockDiagonalPreconditioner<CRDoubleMatrix>*> >
    (prec_pt)->enable_two_level_parallelisation();
```

- **Block Upper Triangular Preconditioning.** Both block triangular preconditioners are implemented in the class `BlockTriangularPreconditioner`. By default this employs the upper-triangular version of the preconditioner.

```
// Block Upper Triangular
prec_pt = new BlockTriangularPreconditioner<CRDoubleMatrix>;
```

- **Block Lower Triangular Preconditioning.** The lower triangular version of the preconditioner can be selected with a call to the method `lower_triangular()`.

```
// Block Lower Triangular
prec_pt = new BlockTriangularPreconditioner<CRDoubleMatrix>;
dynamic_cast<BlockTriangularPreconditioner<CRDoubleMatrix*>*>
(prec_pt)->lower_triangular();
```

Having chosen a preconditioner structure, the next stage is to choose the preconditioner for the subsidiary systems (J_{xx} and J_{yy} in the [Introduction](#)). By default this is `SuperLUPreconditioner`, but we wish to use `HyprePreconditioner` so we pass the previously specified function `Hypre_Subsidary_Preconditioner_Helper::get_hypre_preconditioner()` to the preconditioner.

```
// Specify Hypre as the subsidiary block preconditioner
prec_pt->set_subsidary_preconditioner_function
(Hypre_Subsidary_Preconditioner_Helper::get_hypre_preconditioner);
```

The same subsidiary preconditioner is used for all subsidiary systems in a general purpose block preconditioner.

As discussed in the [Block Preconditioners](#) tutorial, the classification of the DOFs is implemented at an elemental level so we pass a pointer to the mesh containing the elements to the preconditioner. (Note that this problem contains two meshes, one containing the bulk elements and one containing the `FaceElements` that apply the traction boundary condition. Since the latter do not introduce any new DOFs, all the DOFs are classified by the bulk elements. Therefore, we do not need to pass the traction element mesh to the block preconditioner.)

```
// The preconditioner only requires the bulk mesh since its
// elements are capable of classifying all degrees of freedom
// prec_pt is a GeneralPurposeBlockPreconditioner, so we call the function
// add_mesh(...).
prec_pt->add_mesh(problem.solid_mesh_pt());
```

Finally, we pass the preconditioner to the solver

```
// pass the preconditioner to the solver
solver_pt->preconditioner_pt() = prec_pt;
```

and solve the problem:

```
// solve the problem
problem.newton_solve();
```

1.3 Parallelisation

Given that `BlockPreconditioner`, `TrilinosAztecOOSolver`, `SuperLUPreconditioner`, `HyprePreconditioner` and `MatrixVectorProduct` are all automatically distributed, all that is required for a distributed solution is to run the executable under MPI with multiple processes.

1.4 Source files for this tutorial

- The source files for the driver code are in

`demo_drivers/mpi/solvers/`

- The driver code is

`demo_drivers/mpi/solvers/airy_cantilever.cc`

1.5 PDF file

A [pdf version](#) of this document is available.